

# Python for Data 25: Chi-Squared Tests

[back to index \(https://www.kaggle.com/hamelg/python-for-data-analysis-index\)](https://www.kaggle.com/hamelg/python-for-data-analysis-index)

Last lesson we introduced the framework of statistical hypothesis testing and the t-test for investigating differences between numeric variables. In this lesson, we turn our attention to a common statistical test for categorical variables: the chi-squared test.

## Chi-Squared Goodness-Of-Fit Test

In our study of t-tests, we introduced the one-way t-test to check whether a sample mean differs from the an expected (population) mean. The chi-squared goodness-of-fit test is an analog of the one-way t-test for categorical variables: it tests whether the distribution of sample categorical data matches an expected distribution. For example, you could use a chi-squared goodness-of-fit test to check whether the race demographics of members at your church or school match that of the entire U.S. population or whether the computer browser preferences of your friends match those of Internet uses as a whole.

When working with categorical data, the values themselves aren't of much use for statistical testing because categories like "male", "female," and "other" have no mathematical meaning. Tests dealing with categorical variables are based on variable counts instead of the actual value of the variables themselves.

Let's generate some fake demographic data for U.S. and Minnesota and walk through the chi-square goodness of fit test to check whether they are different:

In [1]:

```
import numpy as np
import pandas as pd
import scipy.stats as stats
```

In [2]:

```
national = pd.DataFrame(["white"]*100000 + ["hispanic"]*60000 + \
                        ["black"]*50000 + ["asian"]*15000 + ["other"]*3
                        5000)

minnesota = pd.DataFrame(["white"]*600 + ["hispanic"]*300 + \
                        ["black"]*250 + ["asian"]*75 + ["other"]*150)

national_table = pd.crosstab(index=national[0], columns="count")
minnesota_table = pd.crosstab(index=minnesota[0], columns="count")

print( "National")
print(national_table)
print(" ")
print( "Minnesota")
print(minnesota_table)
```

National

col_0	count
0	
asian	15000
black	50000
hispanic	60000
other	35000
white	100000

Minnesota

col_0	count
0	
asian	75
black	250
hispanic	300
other	150
white	600

Chi-squared tests are based on the so-called chi-squared statistic. You calculate the chi-squared statistic with the following formula:

$$\text{sum}\left(\frac{(\text{observed} - \text{expected})^2}{\text{expected}}\right)$$

In the formula, observed is the actual observed count for each category and expected is the expected count based on the distribution of the population for the corresponding category. Let's calculate the chi-squared statistic for our data to illustrate:

In [3]:

```
observed = minnesota_table

national_ratios = national_table/len(national) # Get population ratios

expected = national_ratios * len(minnesota) # Get expected counts

chi_squared_stat = (((observed-expected)**2)/expected).sum()

print(chi_squared_stat)
```

```
col_0
count    18.194805
dtype: float64
```

*Note: The chi-squared test assumes none of the expected counts are less than 5.*

Similar to the t-test where we compared the t-test statistic to a critical value based on the t-distribution to determine whether the result is significant, in the chi-square test we compare the chi-square test statistic to a critical value based on the [chi-square distribution \(https://en.wikipedia.org/wiki/Chi-squared\\_distribution\)](https://en.wikipedia.org/wiki/Chi-squared_distribution). The scipy library shorthand for the chi-square distribution is chi2. Let's use this knowledge to find the critical value for 95% confidence level and check the p-value of our result:

In [4]:

```
crit = stats.chi2.ppf(q = 0.95, # Find the critical value for 95% confidence*
                      df = 4)   # Df = number of variable categories - 1

print("Critical value")
print(crit)

p_value = 1 - stats.chi2.cdf(x=chi_squared_stat, # Find the p-value
                             df=4)

print("P value")
print(p_value)
```

```
Critical value
9.487729036781154
P value
[0.00113047]
```

Since our chi-squared statistic exceeds the critical value, we'd reject the null hypothesis that the two distributions are the same.

You can carry out a chi-squared goodness-of-fit test automatically using the scipy function `scipy.stats.chisquare()`:

In [5]:

```
stats.chisquare(f_obs= observed, # Array of observed counts
                f_exp= expected) # Array of expected counts
```

Out[5]:

```
Power_divergenceResult(statistic=array([18.19480519]), pvalue=array
([0.00113047]))
```

The test results agree with the values we calculated above.

# Chi-Squared Test of Independence

Independence ([https://en.wikipedia.org/wiki/Independence\\_\(probability\\_theory\)](https://en.wikipedia.org/wiki/Independence_(probability_theory))) is a key concept in probability that describes a situation where knowing the value of one variable tells you nothing about the value of another. For instance, the month you were born probably doesn't tell you anything about which web browser you use, so we'd expect birth month and browser preference to be independent. On the other hand, your month of birth might be related to whether you excelled at sports in school, so month of birth and sports performance might not be independent.

The chi-squared test of independence tests whether two categorical variables are independent. The test of independence is commonly used to determine whether variables like education, political views and other preferences vary based on demographic factors like gender, race and religion. Let's generate some fake voter polling data and perform a test of independence:

In [6]:

```

np.random.seed(10)

# Sample data randomly at fixed probabilities
voter_race = np.random.choice(a= ["asian","black","hispanic","other","white"],
                               p = [0.05, 0.15 ,0.25, 0.05, 0.5],
                               size=1000)

# Sample data randomly at fixed probabilities
voter_party = np.random.choice(a= ["democrat","independent","republican"],
                               p = [0.4, 0.2, 0.4],
                               size=1000)

voters = pd.DataFrame({"race":voter_race,
                       "party":voter_party})

voter_tab = pd.crosstab(voters.race, voters.party, margins = True)

voter_tab.columns = ["democrat","independent","republican","row_totals"]

voter_tab.index = ["asian","black","hispanic","other","white","col_totals"]

observed = voter_tab.iloc[0:5,0:3] # Get table without totals for later use
voter_tab

```

Out[6]:

	democrat	independent	republican	row_totals
asian	21	7	32	60
black	65	25	64	154
hispanic	107	50	94	251
other	15	8	15	38
white	189	96	212	497
col_totals	397	186	417	1000

Note that we did not use the race data to inform our generation of the party data so the variables are independent.

For a test of independence, we use the same chi-squared formula that we used for the goodness-of-fit test. The main difference is we have to calculate the expected counts of each cell in a 2-dimensional table instead of a 1-dimensional table. To get the expected count for a cell, multiply the row total for that cell by the column total for that cell and then divide by the total number of observations. We can quickly get the expected counts for all cells in the table by taking the row totals and column totals of the table, performing an outer product on them with the `np.outer()` function and dividing by the number of observations:

In [7]:

```
expected = np.outer(voter_tab["row_totals"][0:5],
                    voter_tab.loc["col_totals"][0:3]) / 1000

expected = pd.DataFrame(expected)

expected.columns = ["democrat", "independent", "republican"]
expected.index = ["asian", "black", "hispanic", "other", "white"]

expected
```

Out[7]:

	democrat	independent	republican
asian	23.820	11.160	25.020
black	61.138	28.644	64.218
hispanic	99.647	46.686	104.667
other	15.086	7.068	15.846
white	197.309	92.442	207.249

Now we can follow the same steps we took before to calculate the chi-square statistic, the critical value and the p-value:

```
In [8]: chi_squared_stat = (((observed-expected)**2)/expected).sum().sum()

print(chi_squared_stat)
```

7.169321280162059

*Note: We call .sum() twice: once to get the column sums and a second time to add the column sums together, returning the sum of the entire 2D table.*

```
In [9]: crit = stats.chi2.ppf(q = 0.95, # Find the critical value for 95% confidence*
                                df = 8)    # *

print("Critical value")
print(crit)

p_value = 1 - stats.chi2.cdf(x=chi_squared_stat, # Find the p-value
                              df=8)

print("P value")
print(p_value)
```

Critical value  
15.50731305586545  
P value  
0.518479392948842

*Note: The degrees of freedom for a test of independence equals the product of the number of categories in each variable minus 1. In this case we have a 5×3 table so  $df = 4 \times 2 = 8$ .*

As with the goodness-of-fit test, we can use scipy to conduct a test of independence quickly. Use stats.chi2\_contingency() function to conduct a test of independence automatically given a frequency table of observed counts:



In [10]:

```
stats.chi2_contingency(observed= observed)
```

Out[10]:

```
(7.169321280162059, 0.518479392948842, 8, array([[ 23.82 ,  11.16 ,
 25.02 ],
        [ 61.138,  28.644,  64.218],
        [ 99.647,  46.686, 104.667],
        [ 15.086,   7.068,  15.846],
        [197.309,  92.442, 207.249]]))
```

The output shows the chi-square statistic, the p-value and the degrees of freedom followed by the expected counts.

As expected, given the high p-value, the test result does not detect a significant relationship between the variables.

## Wrap Up

Chi-squared tests provide a way to investigate differences in the distributions of categorical variables with the same categories and the dependence between categorical variables. In the next lesson, we'll learn about a third statistical inference test, the analysis of variance, that lets us compare several sample means at the same time.

**Next Lesson: [Python for Data 26: ANOVA](https://www.kaggle.com/hamelg/python-for-data-26-ANOVA)**  
**[. \(https://www.kaggle.com/hamelg/python-for-data-26-ANOVA\)](https://www.kaggle.com/hamelg/python-for-data-26-ANOVA)**

[back to index \(https://www.kaggle.com/hamelg/python-for-data-analysis-index\)](https://www.kaggle.com/hamelg/python-for-data-analysis-index)